

stock_market_analysis

June 15, 2019

```
In [41]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from fastai.tabular import add_datepart
from keras.models import Sequential
from keras.layers import Activation, Dense, Dropout, LSTM
from sklearn import neighbors
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import MinMaxScaler
```

```
In [2]: import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: microsoft_dataset = "../dataset/EOD-MSFT.csv"
tata_dataset = "../dataset/NSE-TATAGLOBAL11.csv"
haineken_dataset = "../dataset/EURONEXT-HEIA.csv"
```

```
In [4]: scaler = MinMaxScaler(feature_range=(0, 1))
```

```
In [5]: df = pd.read_csv(microsoft_dataset)
df.head()
```


```
Out [5]:
```

	Date	Open	High	Low	Close	Volume	Dividend	Split	\
0	2017-12-28	85.90	85.9300	85.550	85.72	10594344.0	0.0	1.0	
1	2017-12-27	85.65	85.9800	85.215	85.71	14678025.0	0.0	1.0	
2	2017-12-26	85.31	85.5346	85.030	85.40	9891237.0	0.0	1.0	
3	2017-12-22	85.40	85.6300	84.920	85.51	14145841.0	0.0	1.0	
4	2017-12-21	86.05	86.1000	85.400	85.50	17990745.0	0.0	1.0	

	Adj_Open	Adj_High	Adj_Low	Adj_Close	Adj_Volume
0	83.768909	83.798164	83.427592	83.593374	10594344.0
1	83.525111	83.846924	83.100903	83.583622	14678025.0
2	83.193546	83.412574	82.920492	83.281313	9891237.0
3	83.281313	83.505607	82.813221	83.388584	14145841.0
4	83.915187	83.963947	83.281313	83.378832	17990745.0

```
In [6]: df["Date"] = pd.to_datetime(df.Date, format="%Y-%m-%d")
df.index = df["Date"]
```

```
plt.figure(figsize=(15, 8))
plt.plot(df["Close"])
plt.xlabel("Date")
plt.ylabel("Final price in particular day")
plt.show()
```



Moving avarage prediction

```
In [7]: def process_dataframe(df):
    data = df.sort_index(ascending=True, axis=0)
    processed_data = pd.DataFrame(index=range(0, len(df)), columns=["Date", "Close"])

    for i in range(0, len(data)):
        processed_data["Date"][i] = data["Date"][i]
        processed_data["Close"][i] = data["Close"][i]

    return processed_data

def plot_predictions(train_data, test_data, predictions):
    test_data.loc[:, "Predictions"] = predictions
    plt.figure(figsize=(15, 8))
    plt.plot(train_data['Close'])
    plt.plot(test_data[['Close', 'Predictions']])
    plt.xlabel("Date")
    plt.ylabel("Final price in particular day")
    plt.show()
```

```

In [8]: lr_data = process_dataframe(df)

In [9]: train_data, test_data = train_test_split(lr_data, test_size=0.2, shuffle=False)

In [10]: def moving_avg(train_data, test_len):
    predictions = []
    train_len = train_data.shape[0]
    for i in range(0, test_len):
        a = train_data["Close"][train_len-test_len+i:].sum() + sum(predictions)
        b = a / test_len
        predictions.append(b)
    return predictions


In [11]: mv_avg_predictions = moving_avg(train_data, test_data.shape[0])
    r_mean_square = np.sqrt(np.mean(np.power((np.array(test_data["Close"]) - mv_avg_predictions)

In [12]: r_mean_square

Out[12]: 15.183666454508193

In [13]: plot_predictions(train_data, test_data, mv_avg_predictions)

```



Moving average approach is easy to implement, but in stock market spectrum predicted data is not looking realistic at all. I will skip linear regression approach. There is a high chance that model will overfit to the date and month column, the model will probably consider the value from the same date a month ago, or the same date/month a year ago.

KNN -> k-nearest neighbours Based on independent variables find similarity between new data and old data. I will use preimplemented KNN from sklearn library.

```
In [14]: df['Date'] = pd.to_datetime(df.Date,format='%Y-%m-%d')
        df.index = df['Date']
```

```
        knn_data = process_dataframe(df)
```

```
In [15]: add_datepart(knn_data, "Date")
        knn_data.drop("Elapsed", axis=1, inplace=True)
```

```
In [16]: train_data, test_data = train_test_split(knn_data, test_size=0.2, shuffle=False)
```

```
In [17]: x_train_scaled = scaler.fit_transform(train_data.drop('Close', axis=1))
        x_test_scaled = scaler.fit_transform(test_data.drop('Close', axis=1))
        x_train = pd.DataFrame(x_train_scaled)
        x_test = pd.DataFrame(x_test_scaled)
        y_train = train_data['Close']
        y_test = test_data['Close']
```

```
In [18]: parameters = {'n_neighbors': [2, 3, 4, 5, 6, 7, 8, 9]}
        knn = neighbors.KNeighborsRegressor()
        model = GridSearchCV(knn, parameters, cv=5)
```


```
In [19]: model.fit(x_train, y_train)
        predictions = model.predict(x_test)
```

```
In [20]: r_mean_square = np.sqrt(np.mean(np.power((np.array(y_test) - np.array(predictions)), 2)))
```

```
In [21]: r_mean_square
```

```
Out[21]: 34.54203483310681
```

```
In [22]: plot_predictions(train_data, test_data, predictions)
```



KNN prediction is not accurate as we can see ://

LSTM - Long Short Term Memory LSTM is artificial recurrent neural network (RNN) architecture used in the field of deep learning. LSTM networks are well-suited to classifying, processing and making predictions based on time series data, since there can be lags of unknown duration between important events in a time series. A common LSTM unit is composed of a cell, an input gate, an output gate and a forget gate. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.

```
In [67]: lstm_data = process_dataframe(df)

        lstm_data.index = lstm_data.Date
        lstm_data.drop('Date', axis=1, inplace=True)

In [68]: train_data, test_data = train_test_split(lstm_data, test_size=0.2, shuffle=False)
        scaled_data = scaler.fit_transform(lstm_data)

In [69]: x_train, y_train = [], []
        for i in range(60, len(train_data)):
            x_train.append(scaled_data[i-60:i,0])
            y_train.append(scaled_data[i,0])
        x_train, y_train = np.array(x_train), np.array(y_train)

        x_train = np.reshape(x_train, (x_train.shape[0],x_train.shape[1],1))

In [70]: class LSTMNet:
        @staticmethod
        def build(height: int):
            input_shape = (height, 1)

            model = Sequential()
            model.add(LSTM(50, return_sequences=True, input_shape=input_shape))
            model.add(LSTM(50))
            model.add(Dense(1))

            return model

In [71]: height = x_train.shape[1]
        model = LSTMNet.build(height)
        model.summary()
```

Layer (type)	Output Shape	Param #
lstm_13 (LSTM)	(None, 60, 50)	10400
lstm_14 (LSTM)	(None, 50)	20200

```
-----
dense_7 (Dense)                (None, 1)                51
=====
```

```
Total params: 30,651
Trainable params: 30,651
Non-trainable params: 0
-----
```

```
In [72]: model.compile(loss='mean_squared_error', optimizer='adam', metrics=['accuracy'])
         model.fit(x_train, y_train, epochs=1, batch_size=1, verbose=2)
```

```
Epoch 1/1
- 86s - loss: 0.0018 - acc: 0.0000e+00
```

```
Out[72]: <keras.callbacks.History at 0x249c1ffccc0>
```

```
In [73]: inputs = lstm_data[len(lstm_data) - len(test_data) - 60:].values
         inputs = inputs.reshape(-1,1)
         inputs = scaler.transform(inputs)
```

```
         X_test = []
         for i in range(60, inputs.shape[0]):
             X_test.append(inputs[i-60:i, 0])
         X_test = np.array(X_test)

         X_test = np.reshape(X_test, (X_test.shape[0],X_test.shape[1],1))
```


```
In [74]: predictions = model.predict(X_test)
         predictions = scaler.inverse_transform(predictions)
```

```
In [75]: r_mean_square=np.sqrt(np.mean(np.power((test_data - predictions), 2)))
```

```
In [76]: r_mean_square
```

```
Out[76]: Close      1.47423
         dtype: float64
```

```
In [77]: plot_predictions(train_data, test_data, predictions)
```



In []: